

Index Number: 190337X

Name: C. J. Kurukulasuriya

```
In [ ]: import tensorflow as tf
        from tensorflow import keras
        from tensorflow.keras import datasets, layers, models
        import numpy as np
        import matplotlib.pyplot as plt
```

Q1) LeNet5 network for MNIST

```
In [ ]: mnist = keras.datasets.mnist
        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

        print("Before padding:")
        print('train_images.shape: ', train_images.shape)
        print('test_images.shape:', test_images.shape)
        print()

        # Padding
        paddings = tf.constant([[0, 0], [2, 2], [2, 2]])
        train_images = tf.pad(train_images, paddings, constant_values=0)
        test_images = tf.pad(test_images, paddings, constant_values=0)

        print("After padding:")
        print('train_images.shape: ', train_images.shape)
        print('test_images.shape:', test_images.shape)
        print('train_labels.shape: ', train_labels.shape)
        print('test_labels.shape:', test_labels.shape)
        class_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']

        #print(train_images.dtype)

        #change dtype
        train_images = tf.dtypes.cast(train_images, tf.float32)
        test_images = tf.dtypes.cast(test_images, tf.float32)
        train_images, test_images = train_images[... , np.newaxis]/255.0,\
                                     test_images[... , np.newaxis]/255.0

        #print(train_images.dtype)
```

Before padding:
train_images.shape: (60000, 28, 28)
test_images.shape: (10000, 28, 28)

After padding:
train_images.shape: (60000, 32, 32)
test_images.shape: (10000, 32, 32)
train_labels.shape: (60000,)
test_labels.shape: (10000,)

```
In [ ]: # LeNet5 for MNIST (as discussed in class)
        model = models.Sequential()
        model.add(layers.Conv2D(6, (5, 5), activation='relu', input_shape=(32, 32, 1)))
        model.add(layers.AveragePooling2D((2, 2)))
        model.add(layers.Conv2D(16, (5, 5), activation='relu'))
        model.add(layers.AveragePooling2D((2, 2)))
        model.add(layers.Flatten())
        model.add(layers.Dense(120, activation='relu'))
        model.add(layers.Dense(84, activation='relu'))
        model.add(layers.Dense(10))

        model.compile(optimizer='adam',
                      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
        print(model.summary())

        model.fit(train_images, train_labels, epochs=5)
        test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=2)
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
conv2d_17 (Conv2D)          (None, 28, 28, 6)          156

average_pooling2d_2 (Averag (None, 14, 14, 6)          0
ePooling2D)

conv2d_18 (Conv2D)          (None, 10, 10, 16)         2416

average_pooling2d_3 (Averag (None, 5, 5, 16)           0
ePooling2D)

flatten_6 (Flatten)         (None, 400)                 0

dense_16 (Dense)            (None, 120)                 48120

dense_17 (Dense)            (None, 84)                  10164

dense_18 (Dense)            (None, 10)                  850

=====
Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0
=====
None
Epoch 1/5
1875/1875 [=====] - 20s 10ms/step - loss: 0.2024 - accuracy: 0.9385
Epoch 2/5
1875/1875 [=====] - 16s 9ms/step - loss: 0.0676 - accuracy: 0.9792
Epoch 3/5
1875/1875 [=====] - 19s 10ms/step - loss: 0.0486 - accuracy: 0.9848
Epoch 4/5
1875/1875 [=====] - 18s 9ms/step - loss: 0.0385 - accuracy: 0.9880
Epoch 5/5
1875/1875 [=====] - 18s 10ms/step - loss: 0.0325 - accuracy: 0.9902
313/313 - 1s - loss: 0.0331 - accuracy: 0.9882 - 1s/epoch - 4ms/step

```

Q2) CNN for CIFAR 10

```

In [ ]: # CIFAR10
(train_images, train_labels), (test_images,
                                test_labels) = datasets.cifar10.load_data()

#print(test_images.dtype)
# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0
#print(test_images.dtype)
print('train_images.shape: ', train_images.shape)
print('train_labels.shape: ', train_labels.shape)
print('test_images.shape:', test_images.shape)
print('test_labels.shape:', test_labels.shape)

class_names = ['airplane', 'automobile', 'bird', 'cat',
               'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

train_images.shape: (50000, 32, 32, 3)
train_labels.shape: (50000, 1)
test_images.shape: (10000, 32, 32, 3)
test_labels.shape: (10000, 1)

```

```

In [ ]: # CNN for CIFAR10
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPool2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
print(model.summary())

model.fit(train_images, train_labels, epochs=5)

```

```
test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=2)
print(test_accuracy)
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_3 (MaxPooling 2D)	(None, 15, 15, 32)	0
conv2d_6 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_4 (MaxPooling 2D)	(None, 6, 6, 64)	0
conv2d_7 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_5 (MaxPooling 2D)	(None, 2, 2, 128)	0
flatten_2 (Flatten)	(None, 512)	0
dense_5 (Dense)	(None, 64)	32832
dense_6 (Dense)	(None, 10)	650

```
=====
Total params: 126,730
Trainable params: 126,730
Non-trainable params: 0
```

```
None
Epoch 1/5
1563/1563 [=====] - 65s 41ms/step - loss: 1.5034 - accuracy: 0.4483
Epoch 2/5
1563/1563 [=====] - 55s 35ms/step - loss: 1.1233 - accuracy: 0.6028
Epoch 3/5
1563/1563 [=====] - 38s 25ms/step - loss: 0.9668 - accuracy: 0.6624
Epoch 4/5
1563/1563 [=====] - 38s 25ms/step - loss: 0.8636 - accuracy: 0.6988
Epoch 5/5
1563/1563 [=====] - 39s 25ms/step - loss: 0.7941 - accuracy: 0.7234
313/313 - 2s - loss: 0.8826 - accuracy: 0.6960 - 2s/epoch - 8ms/step
0.6959999799728394
```

Q3) model_base for MNIST

```
In [ ]: mnist = keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

# Padding
paddings = tf.constant([[0, 0], [2, 2], [2, 2]])
train_images = tf.pad(train_images, paddings, constant_values=0)
test_images = tf.pad(test_images, paddings, constant_values=0)

print('train_images.shape: ', train_images.shape)
print('train_labels.shape: ', train_labels.shape)
print('test_images.shape:', test_images.shape)
print('test_labels.shape:', test_labels.shape)
class_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']

train_images = tf.dtypes.cast(train_images, tf.float32)
test_images = tf.dtypes.cast(test_images, tf.float32)
train_images, test_images = train_images[... ,
                                         np.newaxis]/255.0, test_images[... , np.newaxis]/255.0

train_images.shape: (60000, 32, 32)
train_labels.shape: (60000,)
test_images.shape: (10000, 32, 32)
test_labels.shape: (10000,)
```

```
In [ ]: model_base = models.Sequential()
model_base.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 1)))
model_base.add(layers.MaxPool2D((2, 2)))
model_base.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_base.add(layers.MaxPool2D((2, 2)))
model_base.add(layers.Conv2D(64, (3, 3), activation='relu'))

model_base.add(layers.Flatten())
model_base.add(layers.Dense(64, activation='relu'))
```

```

model_base.add(layers.Dense(10))

model_base.compile(optimizer=keras.optimizers.Adam(),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])
print(model_base.summary())

model_base.fit(train_images, train_labels, epochs=2)
test_loss, test_accuracy = model_base.evaluate(
    test_images, test_labels, verbose=2)

# saving the weights for the Q4 next
model_base.save_weights('saved_weights/')

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_19 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_12 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_20 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_13 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_21 (Conv2D)	(None, 4, 4, 64)	36928
flatten_7 (Flatten)	(None, 1024)	0
dense_19 (Dense)	(None, 64)	65600
dense_20 (Dense)	(None, 10)	650
Total params: 121,994		
Trainable params: 121,994		
Non-trainable params: 0		
Epoch 1/2		
1875/1875 [=====] - 79s 42ms/step - loss: 0.1335 - accuracy: 0.9597		
Epoch 2/2		
1875/1875 [=====] - 87s 46ms/step - loss: 0.0427 - accuracy: 0.9865		
313/313 - 4s - loss: 0.0292 - accuracy: 0.9915 - 4s/epoch - 13ms/step		

Q4) Loading weights to model_lw

```

In [ ]: # The second network with exact same structure as in 3
model_lw = models.Sequential()
model_lw.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 1)))
model_lw.add(layers.MaxPool2D((2, 2)))
model_lw.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_lw.add(layers.MaxPool2D((2, 2)))
model_lw.add(layers.Conv2D(64, (3, 3), activation='relu'))

model_lw.add(layers.Flatten())
model_lw.add(layers.Dense(64, activation='relu'))
model_lw.add(layers.Dense(10))

model_lw.compile(optimizer=keras.optimizers.Adam(),
                loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                metrics=['accuracy'])

print(model_lw.summary())

# Load the saved weights in 3
model_lw.load_weights('saved_weights/')

# train for two epochs
model_lw.fit(train_images, train_labels, epochs=2)
test_loss, test_accuracy = model_lw.evaluate(test_images, test_labels, verbose=2)

# save this model for Q5
model_lw.save('saved_model/')

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
conv2d_22 (Conv2D)          (None, 30, 30, 32)      320

max_pooling2d_14 (MaxPoolin (None, 15, 15, 32)      0
g2D)

conv2d_23 (Conv2D)          (None, 13, 13, 64)      18496

max_pooling2d_15 (MaxPoolin (None, 6, 6, 64)        0
g2D)

conv2d_24 (Conv2D)          (None, 4, 4, 64)        36928

flatten_8 (Flatten)         (None, 1024)             0

dense_21 (Dense)            (None, 64)               65600

dense_22 (Dense)            (None, 10)               650

=====
Total params: 121,994
Trainable params: 121,994
Non-trainable params: 0

```

```

None
Epoch 1/2
1875/1875 [=====] - 87s 46ms/step - loss: 0.0304 - accuracy: 0.9905
Epoch 2/2
1875/1875 [=====] - 73s 39ms/step - loss: 0.0232 - accuracy: 0.9925
313/313 - 5s - loss: 0.0303 - accuracy: 0.9906 - 5s/epoch - 15ms/step
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 3 of 3). These functions will not be directly callable after loading.
INFO:tensorflow:Assets written to: saved_model/assets
INFO:tensorflow:Assets written to: saved_model/assets

```

Q5) Loading model

```

In [ ]: # Loading the above model as model_ld
model_ld = keras.models.load_model('saved_model/')
print(model_ld.summary())
model_ld.evaluate(test_images, test_labels, verbose=2)

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
conv2d_22 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPoolin g2D)	(None, 15, 15, 32)	0
conv2d_23 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPoolin g2D)	(None, 6, 6, 64)	0
conv2d_24 (Conv2D)	(None, 4, 4, 64)	36928
flatten_8 (Flatten)	(None, 1024)	0
dense_21 (Dense)	(None, 64)	65600
dense_22 (Dense)	(None, 10)	650

```

=====
Total params: 121,994
Trainable params: 121,994
Non-trainable params: 0

```

```

None
313/313 - 3s - loss: 0.0303 - accuracy: 0.9906 - 3s/epoch - 10ms/step

```

```

Out[ ]: [0.030290938913822174, 0.990599898910522]

```

Q6) Transfer Learning

```

In [ ]: base_inputs = model_ld.layers[0].input
base_outputs = model_ld.layers[-2].output
output = layers.Dense(10)(base_outputs)

new_model = keras.Model(inputs=base_inputs, outputs=output)

```

```

new_model.compile(optimizer=keras.optimizers.Adam(),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])
print(new_model.summary())

new_model.fit(train_images, train_labels, epochs=3, verbose=2)
new_model.evaluate(test_images, test_labels, verbose=2)

```

Model: "model_3"

Layer (type)	Output Shape	Param #
conv2d_22_input (InputLayer)	(None, 32, 32, 1)	0
conv2d_22 (Conv2D)	(None, 30, 30, 32)	320
max_pooling2d_14 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_23 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_24 (Conv2D)	(None, 4, 4, 64)	36928
flatten_8 (Flatten)	(None, 1024)	0
dense_21 (Dense)	(None, 64)	65600
dense_23 (Dense)	(None, 10)	650
Total params: 121,994		
Trainable params: 121,994		
Non-trainable params: 0		

```

None
Epoch 1/3
1875/1875 - 56s - loss: 0.0820 - accuracy: 0.9773 - 56s/epoch - 30ms/step
Epoch 2/3
1875/1875 - 63s - loss: 0.0192 - accuracy: 0.9938 - 63s/epoch - 33ms/step
Epoch 3/3
1875/1875 - 60s - loss: 0.0145 - accuracy: 0.9955 - 60s/epoch - 32ms/step
313/313 - 4s - loss: 0.0295 - accuracy: 0.9925 - 4s/epoch - 11ms/step

```

Out[]: [0.029508670791983604, 0.9925000071525574]

Q7) Fine Tuning

```

In [ ]: # Loading the saved model
model_for_t1 = keras.models.load_model('saved_model/')
# making the loaded layers non-trainable
model_for_t1.trainable = False
for layer in model_for_t1.layers:
    assert layer.trainable == False

# repeat process in Q6
base_inputs = model_for_t1.layers[0].input
base_outputs = model_for_t1.layers[-2].output
output = layers.Dense(10)(base_outputs)

new_model = keras.Model(inputs=base_inputs, outputs=output)
new_model.compile(optimizer=keras.optimizers.Adam(),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])

new_model.fit(train_images, train_labels, epochs=3, verbose=2)
new_model.evaluate(test_images, test_labels, verbose=2)

```

```

Epoch 1/3
1875/1875 - 18s - loss: 0.2527 - accuracy: 0.9451 - 18s/epoch - 10ms/step
Epoch 2/3
1875/1875 - 18s - loss: 0.0168 - accuracy: 0.9954 - 18s/epoch - 10ms/step
Epoch 3/3
1875/1875 - 20s - loss: 0.0124 - accuracy: 0.9965 - 20s/epoch - 11ms/step
313/313 - 3s - loss: 0.0246 - accuracy: 0.9924 - 3s/epoch - 11ms/step

```

Out[]: [0.024571765214204788, 0.9923999905586243]